



UNIVERSITY OF ALBERTA

DEPARTMENT OF
COMPUTING SCIENCE

An Introduction to Berkeley DB

CMPUT 291

**File and Database Management
Systems**

Objectives

- Levels of abstraction in a typical DBMS.
- Introduction to Berkeley DB.
 - What is Berkeley DB.
 - File Organizations supported by Berkeley DB.
 - Access method operations, cursors.
- Take a look at the Berkeley DB API using simple examples in Python.



Levels of abstraction in a typical DBMS

- Data is stored in tables in conceptual schema.
- Use high level SQL commands (DDL or DML) to manipulate data in database.

Example of DDL: CREATE/DROP TABLE

DML: SELECT...FROM...

- Data is organized in files and indexes in physical schema.
- A DBMS uses suitable storage structure, index files, and access paths to evaluate queries and return results.



What Is Berkeley DB?

- It is an open source embedded database library that provides a simple function-call API for data access and management.
- It supports a number of file structures and operations.
- You can use it to:
 - (1) Develop programs working with storage structures and indexes.
 - (2) Develop applications that don't need the full functionality of a DBMS and require high performance.



What Is Berkeley DB?

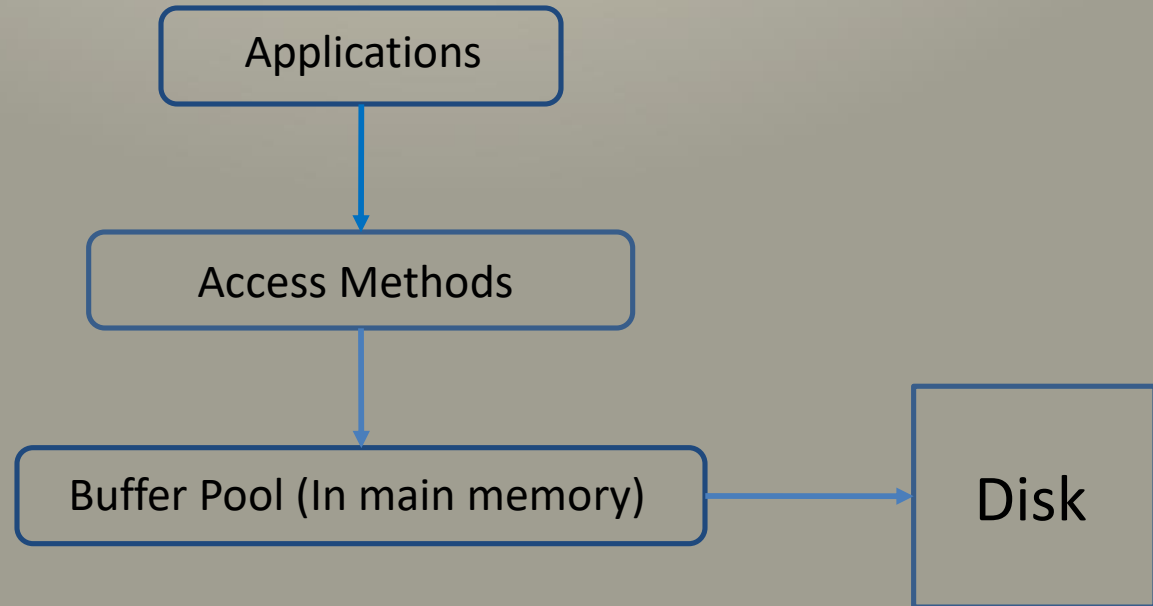


Fig. Abstract Structure of Berkeley DB

The application makes calls into the access methods, and the access methods use the underlying shared memory buffer cache to hold recently used file pages in main memory.

File organizations supported in Berkeley DB

- **Hash:**
 - ✓ Data is stored in an extended linear hash table
 - ✓ Good for applications that need quick random-access.
- **Btree:**
 - ✓ Data is stored in a sorted, balanced tree structure
 - ✓ Good for range-based searches

For more information see [BerkeleyDB Access Methods](#)



File organizations supported in Berkeley DB (Cont.)

Record-number-based:

The logical record number is the primary key of records (compared to Hash and Btree organizations when one or a combination of attributes form a key for records)

- **Queues:**

- ✓ Data is stored in a queue as fixed-length records
- ✓ Good for fast inserts at the tail of the queue
- ✓ It also supports read/delete operation from the head of the queue.
- ✓ provides record-level locking (as opposed to the page-level locking that the other access methods use)

- **Recnos:**

- ✓ Data is stored in either fixed or variable-length records.

For more information see [BerkeleyDB Access Methods](#)



A summary of methods in Database class

Method	Description
open	Create/Open a database
close	Close a database
get	Get items from a database
put	Store items into a database
delete	Delete items from a database
stat	Return database statistics
truncate	Emptying a database of all records
sync	Flushes all modified records from the DB cache to disk
associate	Declares one DB as secondary index of a primary DB

For more information see [BerkeleyDB API](#)



A summary of methods in Cursor class

- A database cursor supports traversing the database.
- It is the only way to access individual duplicate records.

Method	Description
dup	Duplicates a cursor
close	Discards the cursor
count	Returns the number of duplicate data items for the key to which the cursor refers
delete	Deletes the key/value pair to which the cursor refers
get	Retrieves key/value pairs from the database
set	Moves the cursor to the specified key in the database and return the associated key/value pair
put	Stores key/value pairs into the database using a cursor

For more information see [BerkeleyDB API](#)



Berkeley DB API (Python)

1. Create the database:

```
from bsddb3 import db

database = db.DB() #handle for Berkeley DB database
DB_File = "fruit.db"
database.open(DB_File ,None, db.DB_HASH, db.DB_CREATE)
# database.open(DB_File ,None, db.DB_BTREE, db.DB_CREATE)
# database.open(DB_File ,None, db.DB_QUEUE, db.DB_CREATE)
# database.open(DB_File ,None, db.DB_RECNO, db.DB_CREATE)
# The arguments correspond to (fileName, database name within the file for multiple
databases, database type, flag to create database)
```



Berkeley DB API (Python)

2. Declare the cursor for the database:

```
curs = database.cursor()
```

3. Insert records into the database:

```
# insertion of (key, value) pair using the cursor:  
#The arguments corresponds to (key, value, flag to insert key-value pair )  
curs.put(b"apple", "red ", db.DB_KEYFIRST)  
  
# insertion using the database object's put method  
database.put(b"pear", "green")
```



Berkeley DB API (Python)

4. Iterate through the database to display all key-value pairs:

```
# use the cursor
iter = curs.first()
while iter:
    print(iter)
    iter = curs.next()
```



Berkeley DB API (Python)

5. Retrieve a specific key-value or index-value pair

```
# using the database object's get method: only retrieves the value
```

```
result = database.get(b'pear')
```

```
print(result)
```

```
# b'green'
```

```
# using the cursor object's set method:
```

```
result = curs.set(b'pear')
```

```
#Moves the cursor to the specified key in the database and return the associated key/value pair.
```

```
print(result)
```



Berkeley DB API (Python)

6. Remove a (key,value) pair

```
# remove using the cursor – deletes the key-value pair currently referenced by the cursor:  
curs.current() #Returns the key/value pair currently referenced by the cursor  
curs.delete()  
# our database now has (apple,red) only  
  
# remove by using database object  
database.delete(b'apple')  
# Now our database is empty.
```

7. Close database and cursor:

```
curs.close()  
database.close()
```



Berkeley DB API (Python)

8. Show records with the same key (duplicates):

```
# Set duplicate flag before you create the database:
```

```
database.set_flags(db.DB_DUP)
```

```
#Duplicate key-value pairs could be traversed using cursor
```

```
# prints no. of k-v pairs that have the same key (for the key which the cursor is pointing to)  
print(curs.count())
```

```
# prints the next k-v pair if it is a duplicate
```

```
print(curs.next_dup())
```

```
#returns 'None' if the next k-v pair of the database is not a duplicate data record for the  
current key/value pair
```



Exercise 1: Iterate_All

- ❖ Write a code to iterate through all key-value pairs(including the duplicates) using next_dup() method for duplicate keys.
 - In order to test your code, add the following key-value pairs to see if they are retrieved:

```
database.put(b'key1', "value1")
database.put(b'key1', "value2")
database.put(b'key2', "value1")
database.put(b'key2', "value2")
```



Exercise 1: Iterate_All (Cont.)

❖ Iterating through k-v pair including duplicates:

```
curs = database.cursor()
iter = curs.first()

while (iter):

    print(curs.count())
    #prints no. of rows that have the same key for the current key-value pair referred by cursor.

    print(iter)

    #iterating through duplicates:
    dup = curs.next_dup()
    while(dup!=None):
        print(dup)
        dup = curs.next_dup()

    iter = curs.next()
```

❖ A solution is provided on e-class (named Iterate_All.py)



Exercise 2: Populate bdb

Write a program that:

- Gets a student name(as a key) from input
- Searches the database if there is any key(student) with the same name
- Prints any found key as well as its value
- Then, asks the user if the input key should be inserted, and if yes, the program asks for its value as well (let's say the value is the student's mark)
- Terminates if "q" is entered.
 - Pay attention that duplicate keys should be supported (multiple marks for a student)
 - Hint:
 - ✓ Convert Char String to Byte String: `name.encode("utf-8")`
 - ✓ Convert Byte String to Char String: `value.decode("utf-8")`



Exercise 2: Populate bdb (Cont.)

- ❖ In order to set the cursor on the first data item for the input name:

```
name = input("Enter a student Name to look up: ")  
result = curs.set(name.encode("utf-8"))
```

- ❖ To get rid of the byte prefix (b' ') for printing students' names and marks:

```
print("Name: "+str(result[0].decode("utf-8"))+", Mark: "+str(result[1].decode("utf-8")))
```

- ❖ A solution is provided on e-class (named Populate_bdb.py)



Exercise 3: Range Search

Write a program that:

- Gets a name called `Starting_Name`
- Gets another name called `Ending_Name`
- Searches for all students' names from the previous exercise that come after or in the same position as `Starting_Name` (when sorted alphabetically) and come before `Ending_Name`.
- Prints the names and marks of the students found in the previous step.

➤ Hint:

- ✓ In Btree DB type, By default, the sort order is lexicographical, with shorter keys coming before longer keys.
- ✓ `Curs.set_range(key)`: In Btree DB type, sets the cursor to key/value pair which is the smallest key greater than or equal to the specified key, and returns the pair.



Exercise 3: Range Search

- ❖ To get the record that has the smallest key greater than or equal to the Starting Name:

```
Starting_Name = input("Enter the Starting_Name: ")  
result = curs.set_range(Starting_Name.encode("utf-8"))
```

- ❖ To check if the student's name comes before Ending_Name:

```
if (str(result[0].decode("utf-8"))[0:len(Ending_Name)]) < Ending_Name):  
    ...
```

- ❖ A solution is provided on e-class (named Range_Search.py)
- ❖ **Question:** What if the keys were **numbers** instead of names, and you were asked to output the k-v pairs with keys **higher** than **Starting_Number** and **Lower** than **Ending_Number**?



What's Next?

- Check out the page “[introduction to Python bsddb3](#)” for more information.
- Take a look at Berkeley DB API for C & Java in this [tutorial](#).

