

CMPUT 291

File and Database Management Systems

Davood Rafiei
University of Alberta

drafiei@ualberta.ca

Original material Copyright 2001-2019
(some material from textbooks and other instructors)

Fall 2019



Class Information

■ Lectures:

- A1: Mon, Wed & Fri 1:00pm – 1:50pm in ETLC E1007
- A2: Tue & Thu 9:30 – 10:50am in ETLC E1013

■ Instructors:

- Davood Rafiei (drafiei@ualberta.ca)

■ Office hours:

- in ATH 436, time: Wed & Fri 2-3pm

■ TA's:

- Names and contacts in the course web page



Labs

- **Location:** CSC 219

(CSC B10 and ETLC E2002 for the laptop labs)

- **10 labs:** schedule posted in the course web page

- **Working with:**

- Unix tools and commands
- SQLite
- Berkeley DB

- Access from home

- Assignments and projects are tested on lab machines



Textbook and Slides

- **Recommended Text:**

- Silberschatz, Korth, Sudarshan: Database Systems Concepts, 7th edition, McGraw Hill, 2019.

- **Alternative Text:**

- Kifer, Bernstein, Lewis;
Database Systems : An Application Oriented
Approach , 2nd edition, Addison Wesley, 2005


- **Schedule, slides, readings and forums:**

- Check the eclass pages at
<https://eclass.srv.ualberta.ca/course/view.php?id=53330>



Evaluation, etc.

- **Assignment 1:** 5% (groups of 2-3)
 - Available 2 weeks in advance
- **Assignment 2:** 10% (**individual**)
 - Available 2 weeks in advance
- **Quizzes:** 7% total
 - 2 in the lab (4%), 3 in lectures (3%)
- **Mini-Project 1:** 13% (groups of 2-3)
 - Available 3 weeks in advance
- **Mini-Project 2:** 12% (groups of 2-3)
 - Available 2-3 weeks in advance
- **Term Test:** 23%
- **Final Exam:** 30%



This course

- A balanced mix of theory and practice
 - Theory
 - ✓ query languages with roots in logic
 - ✓ set theory
 - ✓ normal forms
 - Practice
 - ✓ writing queries
 - ✓ building databases
 - ✓ developing database applications
 - ✓ searching files and indexes



What is expected?

- A good understanding of the concepts covered
- Hands on experience
 - Writing queries, developing database applications



How to succeed?

- For theory
 - Regularly attend the lectures/labs and ask questions
 - Closely follow the course and make use of office hours to resolve problems
- For practice
 - Write as many queries as possible (like a musician, an artist, a programmer, practice is the key!)
 - Attend the labs and make the best use of it
 - Start the assignments and projects early
- Stay on top
 - Topics covered can be overwhelming




Course Outline

- What are they and why are they useful?
 - separate data from programs.
 - factor out common pieces from applications.
- How can I use them?
 - Structuring data: relational data model
 - queries: SQL
- How do they work?
 - Organizing data in files and indexes
 - Query execution
 - Performance tuning



What is a Database?

- A database is a very large, integrated collection of data.
- Models real-world enterprise.
 - Entities (e.g. students, courses)
 - Relationships (e.g. Justin Trudeau is taking 291).



Manufacturing

Product data

University

Student data, courses

Hospital

Patient data, facilities

Bank

Account data

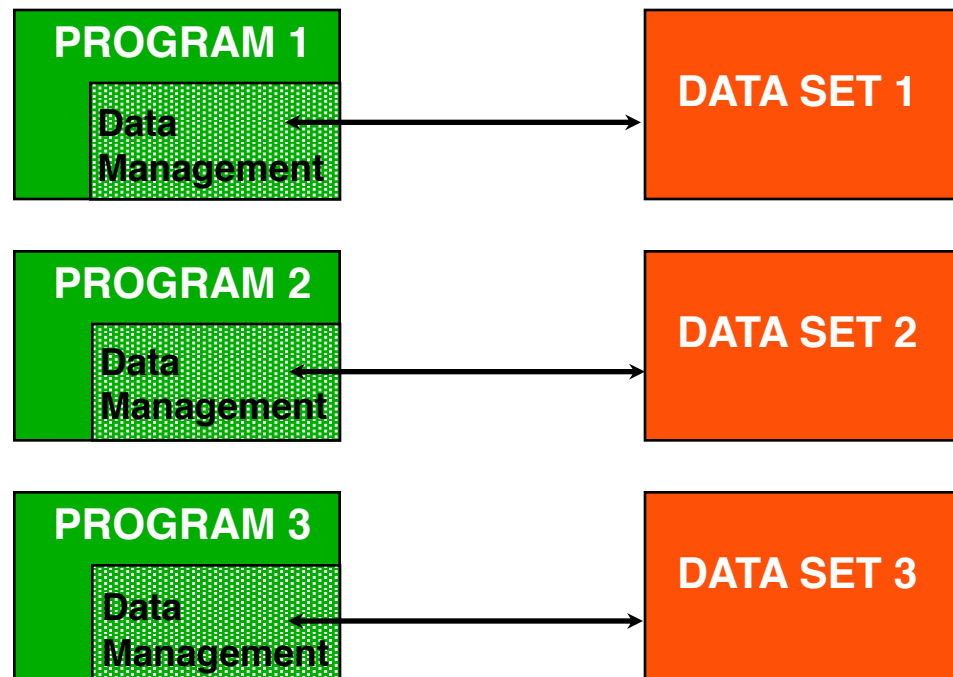
What is a DBMS?

- A Database Management System (DBMS) is a software package designed to store and manage databases.



Non-Database Approach

- Common in 60's.
- One data set per program.
- Programmer defines (and implements) storage structures, access methods, etc.

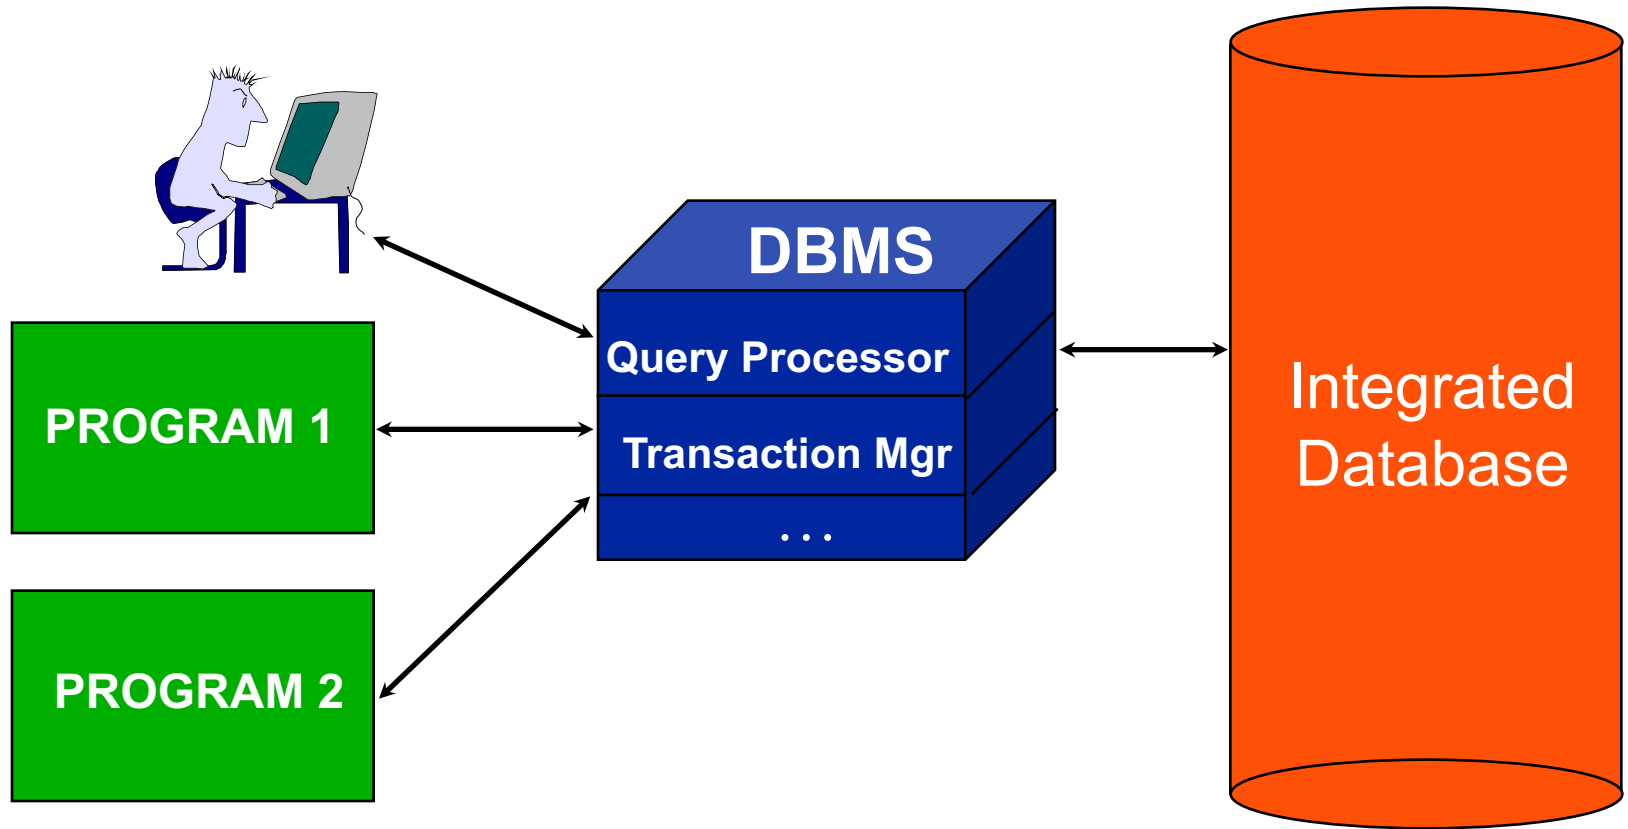


Problems

- With more and more applications we get
 - many files with different structures
 - redundant storage
 - inconsistent copies
 - expensive updates
 - incorrect data
 - data exchange between applications



Database Approach



Basic Idea

- Remove details related to data storage and access from application programs.
- Concentrate those functions in a single system called *database management system*.
- Have all applications access data through the DBMS.



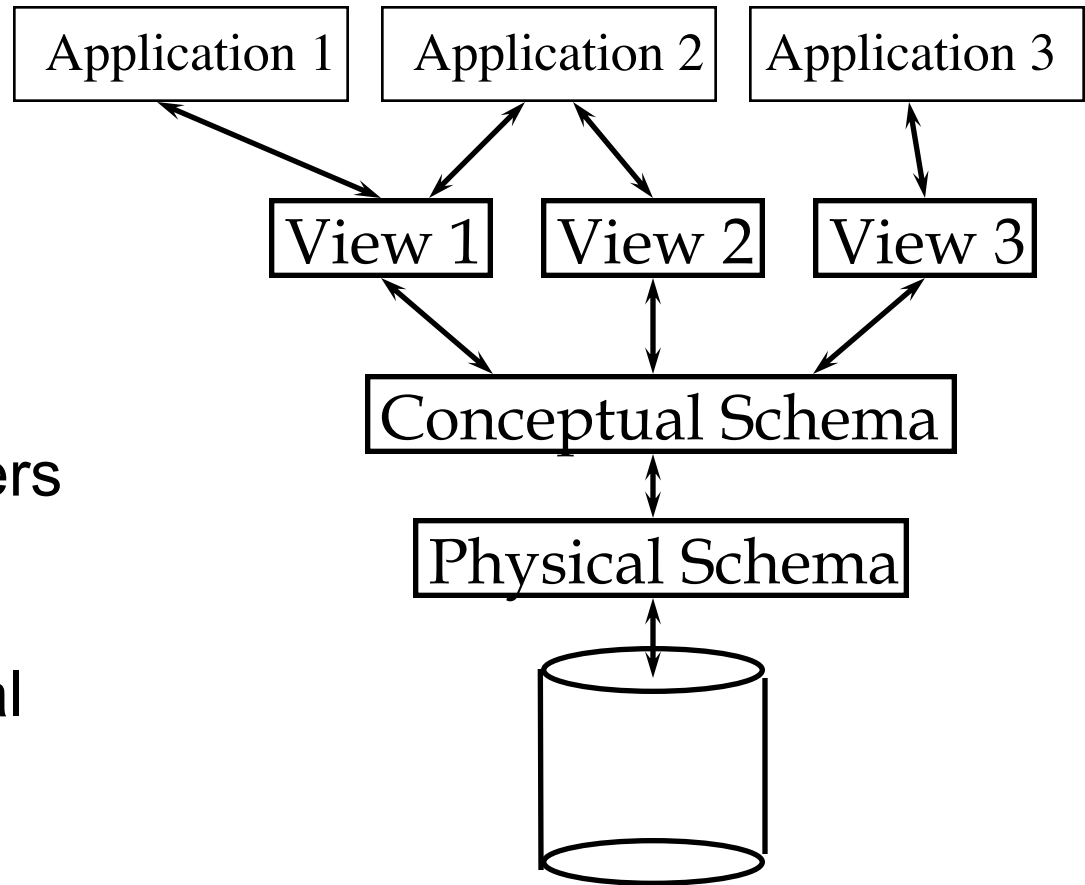
Why Use a DBMS?



- Reduced redundancy
- Less risk of inconsistency
- Reduced application development time
- Uniform data administration
- Concurrent access, recovery from crashes
- **But**, most importantly
Data Independence

Three Schema Levels

- **Schema:** a description of the data contents, structures and other aspects.
- **External schema:** what appl. programs and users see.
- **Conceptual schema:** description of the logical structure of data.
- **Physical schema:** file structures and indexes being used.



Example: University Database

- External Schema (View):
 - *Course_info(cid:char(8),enrollment:int)*
- Conceptual schema:
 - *Students(sid: char(8), name: char(16), login: char(8), age: int, gpa:float)*
 - *Courses(cid: char(8), cname:char(16), credits:int)*
 - *Enrolled(sid:char(8), cid:char(8), grade:char(2))*
- Physical schema:
 - Students is stored as an unordered file with file name data.txt.
 - There is an index on the first column of Students.



Data Independence

- **Objective:** Application programs unaffected by changes in storage structure and access strategy.
- Logical data independence: Protection from changes in *logical* structure of data.
- Physical data independence: Protection from changes in *physical* structure of data.

➡ *One of the most important benefits of using a DBMS!*

DBMS Functionality

- **Data definition** facilities
 - provides a data definition language (DDL),
 - stores the definitions in a user-accessible catalog (data dictionary).
- **Data manipulation** facilities
 - provides a query language for storing, retrieving and updating data.
- Facilities for **integrity constraints**
 - does the validation check for integrity constraints before updates,
 - different kinds of constraints
 - ✓ primary key constraints (entity integrity),
 - ✓ foreign key constraints (no dangling references),
 - ✓ check constraints.

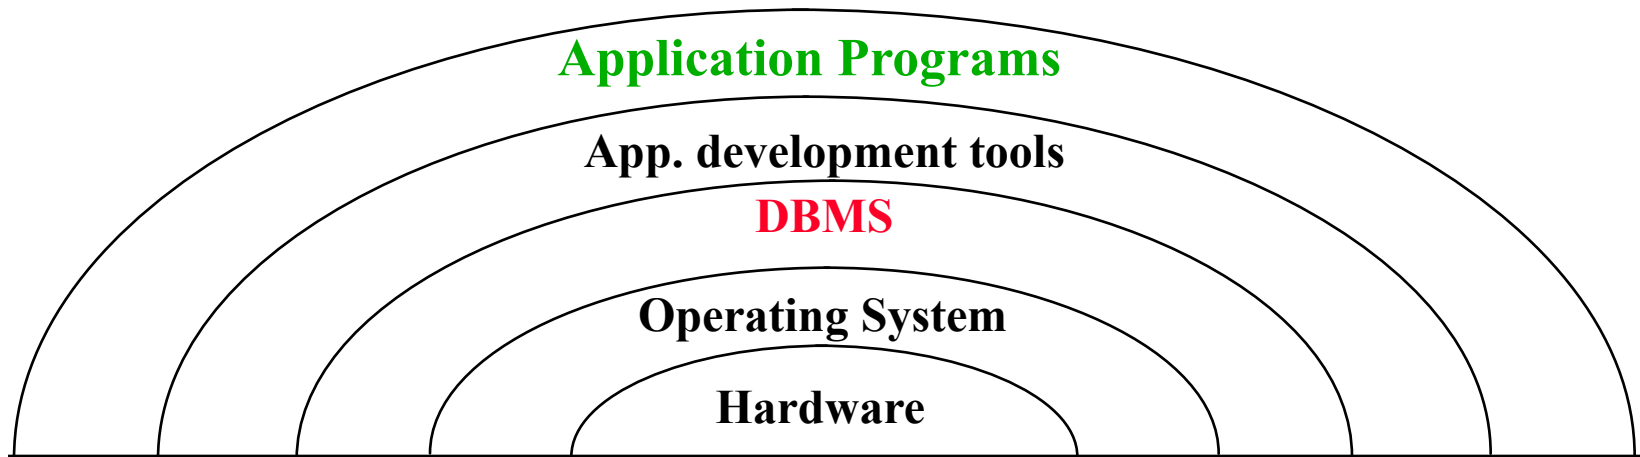


DBMS Functionality (cont'd)

- Provides **concurrency control**
 - multiple users simultaneously access/update a database.
- Supports **transactions**
 - a sequence of user operations to be performed as an atomic action,
 - all operations or none are performed.
- Provides **database recovery**
 - whatever happens, never lose data.
- Provides **query optimization**
 - find the best possible plan for executing a query.



Place in a Computer System



Why this Course?



- A paradigm shift from computation to information
- “Big data” is everywhere!
 - Digital libraries, online stores, the Web, etc.
- Many applications need DBMS functionalities
- One size doesn’t fit all
- Challenge: manage this diversity and volume (need more people).
- It is also fun!

Course Outline

- What are they and why are they useful?
 - Introduction
- How can I use them?
 - Relational model
 - ER model
 - Relational query languages
 - Database design
- How do they work?
 - Secondary storage
 - Organizing data in files and indexes
 - Physical database design
 - Database tuning



Summary

- Introduced databases and DBMS
- Reviewed some functionalities of a DBMS and benefits
- Three-layered schema architecture
- Not sold yet?
 - Google returns 122 million hits for “jobs DBA” compared to 171 million for “jobs programmer”,
 - and good DBAs are **well-paid!**
 - As of Aug 2019, DBAs and DB developers are paid \$77k-\$167k
 - <http://www.itcareerfinder.com/brain-food/it-salaries/database-administrator-salary-range.html>
 - DBMS R&D is one of the broadest and exciting areas in CS

