

SQL Inside Applications

Davood Rafiei



Introduction

- **Basic Idea:** Use SQL statements inside a host language (Java, Python, C/C++, ...).
- **Advantages:**
 - Can do all the fancy things you do in C/Java/Python.
 - Still have the power of SQL.
- **Need mechanisms for**
 - Embedding or calling SQL statements in code
 - Transferring data to/from DBMS
 - Compiling, linking, etc.



SQL inside Applications

- Statement-level interface
 - SQL statements appear as new statements in the program.
 - Precompile step: replaces new statements with some procedure calls.
 - Flavors: **Static SQL**, **Dynamic SQL**
- Call-level interface
 - Program is written entirely in the host language (and can be compiled or executed)
 - SQL statements are passed as parameters to some functions.
 - Flavors: **ODBC**, **JDBC**, **sqlite callback**



Statement-level Interface

- New statements:

- EXEC SQL <sql statement>
- E.g.
 - ✓ EXEC SQL SELECT ...
 - ✓ EXEC SQL UPDATE ...

- Static SQL: e.g.

```
EXEC SQL SELECT COUNT(*) INTO :cnt
        FROM emp;
```

The SQL statement is known at compile time.

- Dynamic SQL: e.g.

```
strcpy(qstr, "SELECT COUNT(*) FROM emp");
EXEC SQL PREPARE Q FROM :qstr;
```

The SQL statement is not known at compile time.



Program Structure

SQL Include statements

```
main(int argc, char *argv[])
```

```
{
```

Declarations



Connect to the Database



Do your work with the database



Process Errors



Commit/Rollback

```
}
```



Host Variables

- Pass values between a SQL statement and the rest of the program.
- Declaration: as follows
EXEC SQL BEGIN DECLARE SECTION;
int cnt;
char name[20];
EXEC SQL END DECLARE SECTION;
- Usage in SQL statements:
EXEC SQL SELECT COUNT() INTO :cnt FROM emp;*

Colon distinguishes host variables from SQL identifiers



Error Handling

- Check `sqlca.sqlcode`, the return code of SQL statements
 - `== 0` : the command was successful
 - `> 0` : no row in the output
 - `< 0` : error



myprog.pc: Our First Program

```
#include <stdio.h>
EXEC SQL INCLUDE SQLCA;

int main(int argc, char *argv[])
{
    EXEC SQL BEGIN DECLARE SECTION;
        char user[10] = "SCOTT";
        char pwd[10] = "TIGER";
    EXEC SQL END DECLARE SECTION;

    EXEC SQL WHENEVER SQLERROR GOTO error;
    EXEC SQL CONNECT :user IDENTIFIED BY :pwd;
    printf("connected to Oracle");

    /* SQL statements */

    EXEC SQL COMMIT RELEASE;
    return 0;
```



myprog.pc (Cont.)

error:

```
sqlca.sqlerrm.sqlerrmc[sqlca.sqlerrm.sqlerrml] = '\0';  
printf("Oracle Error: %s\n", sqlca.sqlerrm.sqlerrmc);
```

```
EXEC SQL WHENEVER SQLERROR CONTINUE;  
EXEC SQL ROLLBACK RELEASE;
```

```
return 1;
```

```
}
```



Program Preparation

- Precompile

 - proc myprog.pc**

 - result: myprog.c (a pure C program with SQL statements replaced with library calls)

- Compile the C program

 - cc myprog.c**

 - result: myprog.o

- Link the libraries

 - cc -o myprog myprog.o -L...**

 - result: myprog (an executable program)

- **** Use a makefile instead ****

 - Check the Oracle directory for a sample *makefile*



myprog2.pc: Our 2nd Program

```
/* get the emp_number from input and print the emp_name */
```

```
...
```

```
EXEC SQL BEGIN DECLARE SECTION;
```

```
    int emp_number;
```

```
    char emp_name[30];
```

```
EXEC SQL END DECLARE SECTION;
```

```
EXEC SQL WHENEVER SQLERROR    GOTO error;
```

```
EXEC SQL WHENEVER NOT FOUND    GOTO nope;
```

```
/* Connect to the database ..... */
```

```
printf("Enter emp_number:");
```

```
scanf("%d", emp_number);
```

```
EXEC SQL SELECT ename INTO :emp_name
```

```
        FROM emp
```

```
        WHERE empno = :emp_number;
```

```
printf("Employee name is %s\n", emp_name);
```

```
return 0;
```

```
error: ...
```

```
nope: ...
```



Complication

- What if a SQL statement returns more than one row?
 - cannot fit it in any C variable!
 - solution: use a **cursor**.



Use of a Cursor

```
/* print the names of all employees */  
...  
EXEC SQL BEGIN DECLARE SECTION;  
    char emp_name[30];  
EXEC SQL END DECLARE SECTION;  
  
EXEC SQL DECLARE emp_cursor CURSOR FOR  
        SELECT ename  
        FROM   emp;  
  
EXEC SQL OPEN emp_cursor;  
EXEC SQL WHENEVER NOT FOUND GOTO end;  
  
printf("Employee names are:\n");  
for (;;) {  
    EXEC SQL FETCH emp_cursor INTO :emp_name;  
    printf("%s\n", emp_name);  
}  
end:  
EXEC SQL CLOSE emp_cursor;  
EXEC SQL COMMIT RELEASE;  
return 0;
```

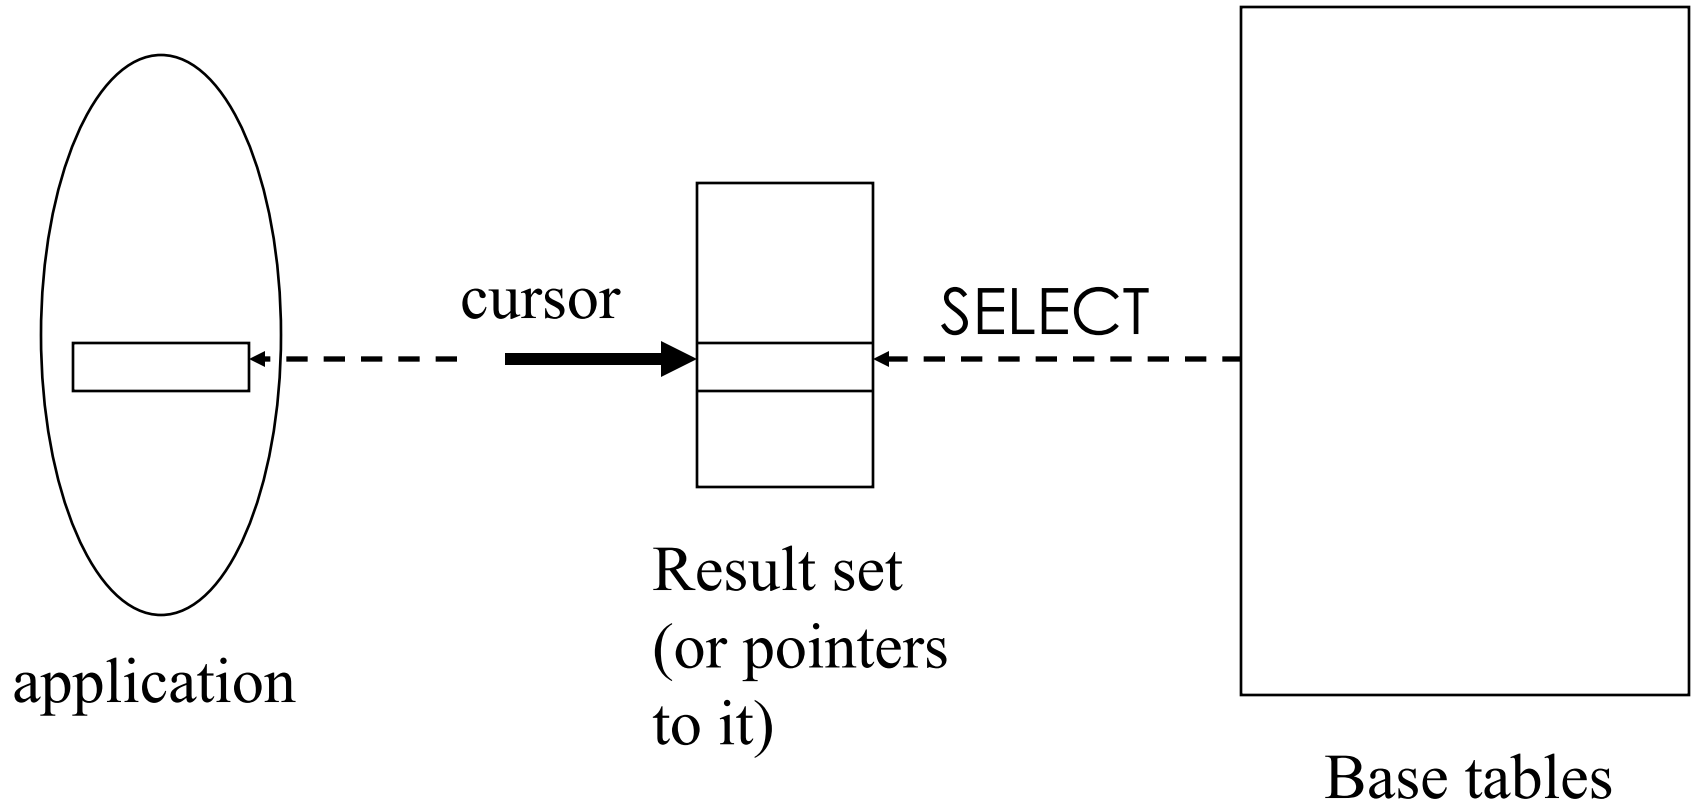


Cursor Overview

- ***Result set*** – set of rows produced by a SELECT statement
- ***Cursor*** – pointer to a row in the result set.
- Cursor operations:
 - *Declaration*
 - *Open* – execute SELECT to determine result set and initialize pointer
 - *Fetch* – advance pointer and retrieve next row
 - *Close* – deallocate cursor



Cursor




Statement-level Interface (Summary)

- One database at a time
- Both schema (for input and output) and database name must be known at compile time
- Ordered from the most efficient to the least
 - Static SQL
 - Dynamic SQL
 - Call-level interface: JDBC



Call-level Interface

- Neither the schema nor the database is known at compile time
- The application can connect to more than one database



- We discuss JDBC and Callback for accessing our databases in Java and C

JDBC

- A Java API (a set of function calls) to communicate with SQL engines
- Supports database queries and updates as well as metadata retrievals (e.g. names of tables and columns)



Steps

- Open a connection

```
Connection conn = DriverManager.getConnection(url);
```

- Create a statement

```
Statement stmt = conn.createStatement();
```

- Execute the queries and fetch the results

```
ResultSet rs = stmt.executeQuery("select id, name, phone from emp;")  
// process the result, one row at a time
```

- Handle errors (exception handling)

```
try {...}  
catch (SQLException e) {  
    System.out.println("SQLException : " + e);  
}
```



A Simple Program (P1.java)

```
import java.sql.*;

public class P1 {

    public static void prog1 () {

        String driverName = "org.sqlite.JDBC";
        String url ="jdbc:sqlite:/Users/drafiei/Courses/291/code/jdbc/291dbfile.db";
        Connection con = null;
        Statement stmt = null;
        ResultSet rs = null;

        try {
            Class.forName(driverName);
        } catch (ClassNotFoundException e) {
            System.out.println("ClassNotFoundException : " + e);
        }
    }
}
```



```

try {
    con = DriverManager.getConnection(url);
    stmt = con.createStatement();
    rs = stmt.executeQuery("select id, name, phone from emp");

    while(rs.next()) {
        System.out.println(rs.getInt(1) + "  " +
            rs.getString(2) + "  " +
            rs.getString(3));
    }
    stmt.close();
    con.close();
}
catch (SQLException sqle) {
    System.out.println("SQLException : " + sqle);
}

}

public static void main (String[] args) { prog1();}
}

```



Compile and Run

- Create table emp(id, name, phone) with some tuples and store it in 291dbfile.db
- Have sqlite-jdbc in class path
- Compile and run

```
javac -cp "../sqlite-jdbc-3.27.2.1.jar" P1.java  
java -cp "../sqlite-jdbc-3.27.2.1.jar:." P1
```

Include the current directory that has the class file



Executing a Query

```
import java.sql.*;    // imports all classes in package java.sql
String driverName = "org.sqlite.JDBC";
String db_url = "jdbc:sqlite:/Users/drafiei/291dbfile.db";
```

```
//String driverName = "oracle.jdbc.driver.OracleDriver";
//String db_url = "jdbc:oracle:thin:@gwynne.cs.ualberta.ca:1521:CRS";
```

```
Class.forName(driverName);    // loads the specified driver
```

```
Connection con = DriverManager.getConnection(db_url)
//Connection con = DriverManager.getConnection(db_url, Id, Pwd);
```

- connects to the DBMS at address *db_url*
- If successful, creates a connection object, *con*, for managing the connection

```
Statement stmt = con.createStatement ();
```

- Creates a statement object *stmt*
- Statements have `executeQuery()` method



Executing a Query

```
String query = "SELECT T.StudId FROM Transcript T" +  
               "WHERE T.CrsCode = 'cse305' " +  
               "AND T.Semester = 'S2000' ";
```

```
ResultSet rs = stmt.executeQuery(query);
```

- *Creates a result set object, rs.*
- *Prepares and executes the query.*
- *Stores the result set produced by execution in rs (analogous to opening a cursor).*
- *The query string can be constructed at run time (as above).*
- *The input parameters are plugged into the query when the string is formed (as above)*



Result Sets and Cursors

- Three types of result sets in JDBC:
 - *Forward-only*: not scrollable
 - *Scroll-insensitive*: scrollable (can jump up and down in the result set!), changes made to underlying tables after the creation of the result set are not visible through that result set
 - *Scroll-sensitive*: scrollable, changes made to the tuples in a result set after the creation of that set are visible through the result set



Scrollable Result Set

```
stmt = con.createStatement(ResultSet.TYPE_SCROLL_INSENSITIVE,  
                           ResultSet.CONCUR_READ_ONLY);  
rs = stmt.executeQuery("select id, name, phone from emp");
```

// playing with a scrolable Result Set

```
rs.absolute(3); // jump to Row 3
```

// Traverse the resultset from bottom

```
rs.afterLast();  
while(rs.previous()) {  
    System.out.println(rs.getInt(1) + " " +  
        rs.getString(2) + " " +  
        rs.getString(3));  
}
```

ResultSet Type Values

- TYPE_FORWARD_ONLY
- TYPE_SCROLL_INSENSITIVE
- TYPE_SCROLL_SENSITIVE

SQLite only supports TYPE_FORWARD_ONLY



Updateable Result Set

```
Statement stat = con.createStatement (ResultSet.TYPE_SCROLL_SENSITIVE,  
                                       ResultSet.CONCUR_UPDATABLE );
```

- Any result set type can be declared *read-only* (**CONCUR_READ_ONLY**) or *updatable* (**CONCUR_UPDATABLE**), assuming SQL query satisfies the conditions for updatable views
- Current row of an updatable result set can be updated/deleted and a new row can be inserted, causing changes in base table

```
rs.updateString ("Name", "John" );    // update attribute "Name" of  
                                       // current row in row buffer.  
or < rs.updateRow ( );                // make the change permanent  
    rs.CancelRowUpdate ( );           // cancel the update
```

SQLite only supports **CONCUR_READ_ONLY**



Accessing Metadata

- Metadata includes table names, column counts, column names and types, etc.

```
rs = stmt.executeQuery("select * from emp");
```

```
// Access metadata
```

```
ResultSetMetaData rsmd = rs.getMetaData();
```

```
for (int i = 1; i <= rsmd.getColumnCount(); i++) {  
    System.out.println(rsmd.getColumnName(i) + " : " +  
        rsmd.getColumnTypeName(i));  
}
```



Controlling Transactions

- Each SQL statement is treated as a transaction and is committed automatically
 - Not always a good strategy (?)
- Alternative
 - Turn off auto commit and commit/rollback explicitly

```
con = DriverManager.getConnection(url);  
con.setAutoCommit(false);
```

```
String sql = "INSERT INTO emp(id,name,phone) VALUES(?,?,?)";  
PreparedStatement pstmt = con.prepareStatement(sql);  
pstmt.setInt(1, 0); pstmt.setString(2, "Bob"); pstmt.setString(3, "780-000-0000");  
pstmt.executeUpdate();
```

```
con.commit(); // con.rollback();
```



JDBC (Summary)

- More flexible than Statement Level Interface
- Schema and table names may not be known in advance
- Scrollable and updatable result set
(though sqlite is limited on this aspect)
- Can connect to any number of databases



SQLite callback in C

- Three functions: open, exec, close
 - `sqlite3_open(const char *filename, sqlite3 **db)`
 - `sqlite3_exec(sqlite3 *db, const char *sql, sqlite_callback, void *data, char **errmsg)`
 - `sqlite3_close(sqlite3 *db)`
- Output is processed by a callback function



myp.c

```
##include <stdio.h>
#include <stdlib.h>
#include <sqlite3.h>

static int callback(void *data, int argc, char **argv, char **aColName){
    int i;
    fprintf(stderr, "%s: ", (const char*)data);

    for(i = 0; i<argc; i++){
        printf("%s = %s\n", aColName[i], argv[i] ? argv[i] : "NULL");
    }

    printf("\n");
    return 0;
}
```




```

int main(int argc, char* argv[]) {
    sqlite3 *db;  char *zErrMsg = 0;
    int rc;  char *sql;
    const char* data = "Callback function called";

    /* Open database */
    rc = sqlite3_open("291lect.db", &db);
    if( rc ) {
        fprintf(stderr, "Can't open database: %s\n", sqlite3_errmsg(db));
        return(0);
    } else {
        fprintf(stderr, "Opened database successfully\n");
    }

    /* Create and execute SQL statement */
    sql = "SELECT * from customer";
    rc = sqlite3_exec(db, sql, callback, (void*)data, &zErrMsg);

    if( rc != SQLITE_OK ) {
        fprintf(stderr, "SQL error: %s\n", zErrMsg);
        sqlite3_free(zErrMsg);
    } else {
        fprintf(stdout, "Operation done successfully\n");
    }
    sqlite3_close(db);
    return 0;
}

```



Compile and Run

- Compile

```
gcc -Wall -std=c99 -L/usr/lib/sqlite3 myp.c -lsqlite3 -o myp
```

- Run

```
./myp
```

Opened database successfully

Callback function called: cname = Davood

street = 114 St

city = Edmonton

Callback function called: cname = Ehsan

street = University Ave

city = NULL

Operation done successfully

More information
https://www.sqlite.org/c_interface.html



More Information

- Statement-level Interface:
 - Not supported in SQLite
 - But is supported in other databases such as Oracle
- Call-level Interface
 - JDBC tutorials
 - Callback interface in C



Summary

- Covered:
 - Static SQL (no SQLite support)
 - Call level interface (JDBC, C)
- SQLite in Python is covered in the lab
- Not Covered: Dynamic SQL, ODBC
- Final note (but quite important):
 - Avoid data processing in the host language if it can be passed to SQL.
 - Use the host language for things that cannot be done in SQL.

